

Successive h-indices using OpenAlex

Nicholas S. Selby ^{1*}

^{1*}Renewvia Solar Ethiopia PLC, Bole sub-city, Woreda 03, House No. 174/175, Saya Building, 5th Floor, Addis Ababa, Ethiopia.

Corresponding author(s). E-mail(s): nicholas.selby@renewvia.com;

Abstract

I apply Schubert’s successive h-index framework to OpenAlex, computing h_2 for 20,932 institutions and h_3 for 214 countries across 30.0 million researchers. The overall institution-level ranking is, in substance, a ranking of medical school size: Medicine accounts for 25% of the author pool and structurally dominates any university-wide aggregate. Field-specific decomposition is more informative: according to this metric, Wageningen University leads Agricultural & Biological Sciences and Environmental Science while ranking 172nd overall; Carnegie Mellon leads Computer Science from 175th; Brazil holds two of the world’s six deepest dental-research benches. I normalize for group size using Egghe’s Lotkaian scaling, which accounts for 91–92% of variance in both h_2 and h_3 , and find that Italy and Japan reach nearly the same h_3 from a $6\times$ difference in institutional scale.

Keywords: successive h-index, OpenAlex, institutional ranking, research assessment

JEL Classification: I23 , C80

MSC Classification: 62P99 , 62-07 , 62P25

1 Introduction

In 2007, András Schubert published a paper titled “Successive h-indices.” (2007) It begins:

Hirsch’s excessively discussed h-index generated a long row of interpretations, applications and derivatives within a surprisingly short period. Symptomatically, Hirsch’s paper in the *Proceedings of the National Academy of Sciences of the United States of America* became the second highest cited

paper of the issue in which it had been published. In this note, attention is called to an apparently so far overlooked option: h-indices themselves may form the basis of a successive series of h-indices.

Given, e.g., the h-indices of the researchers of an institute (which indices will be denoted by h_1 in what follows), an index, h_2 , of the institute can be determined: the institute has an index h_2 if h_2 of its N researchers have an h_1 -index of at least h_2 each, and the other $N - h_2$ researchers have h_1 -indices lower than h_2 each. The succession can then be continued, e.g., for networks of institutions or countries or other higher levels of aggregation.

Since no sufficient data for testing the concept on the researcher-institute-country hierarchy were available, another example was chosen.

Fifteen years later, OpenAlex (Priem, Piwowar, & Orr, 2022) launched a FOSS bibliographic catalogue complete with citation indexing, authors, fields, and institutions. Recent evaluations have found it to be a broadly suitable substitute for proprietary sources such as Scopus and Web of Science at the institutional level (Alperin, Portenoy, Demes, Larivière, & Haustein, 2024; Culbert et al., 2025), and I take it here as the dataset for which Schubert wished in 2007.

The idea of applying the h-index hierarchically at the researcher-institution level was proposed independently by Prathap (2006) in a one-page note less than a month before Schubert’s paper. Schubert’s distinctive contribution was to recognize that the construction extends to an arbitrary succession of levels, opening the path to h_3 and beyond. In the years between, successive h-indices attracted modest but sustained attention. Egghe (2008) gave the concept a theoretical footing, showing how the exponent of the underlying Lotka distribution propagates through each successive aggregation step. Several case studies followed shortly after: Ruane and Tol (2008) computed a rational (non-integer) successive h-index across economics departments in Ireland; Arencibia-Jorge, Barrios-Almaguer, Fernández-Hernández, and Carvajal-Espino (2008) applied Schubert’s exact researcher-department-institution hierarchy to a Cuban research institute using Web of Science data; and Arencibia-Jorge and Rousseau (2009) applied Prathap’s framing to twelve Cuban brain-research institutions. None of this prior work, however, has had access to a comprehensive, open, institution-labeled snapshot of the kind OpenAlex now provides.

Schubert suggested that h_2 of institutes would be more interesting than that of journals; Section 2 describes how I compute it at scale, and Section 3 presents the results.

2 Methodology

2.1 Data source and filtering

The underlying data is OpenAlex’s public authors snapshot (Priem et al., 2022), distributed as Parquet on `s3://openalex/data/parquet/authors/` (`us-east-1`, no credentials required) and partitioned into over 1,600 daily `updated_date=` directories totalling about 53 GB. Only four columns are read from each remote file (`id`, `h_index`, `affiliations`, `topics`), and an author is retained only if all of the following hold:

- `summary_stats.h_index` is non-null and greater than zero;
- at least one entry in `affiliations` has `institution.type = 'education'`;
- at least one entry in `topics` carries a non-null field classification.

Applying these filters to the full snapshot yields 30.0 million authors, spanning 20,932 institutions and OpenAlex’s current list of 26 fields. One limitation carried over from the source data: an author’s `h_index` is computed by OpenAlex from all works meeting its citation criteria regardless of retraction status, which OpenAlex records inconsistently (Hauschke & Nazarovets, 2025); a small, unquantified fraction of h_1 values in the dataset may therefore be inflated by since-retracted work.

2.2 Author-level assignment

Each retained author is reduced to a single row: one institution and one field. An author may list several affiliations with `institution.type = 'education'`; rather than fractionally splitting or duplicating that author across institutions, I assign them to the one with which they were most recently affiliated. Recency is read from OpenAlex’s `affiliations` field, which records, for each institution an author has ever listed, the calendar years in which they published while there; I take the institution whose most recent such year is latest, breaking ties (identical latest year) by the numerically smallest OpenAlex institution identifier. I deliberately do not use the similarly-named `last_known_institutions` field for this: despite its name, it is simply the set of affiliations listed on an author’s single most recent work, carrying no per-entry recency of its own, so when an author has more than one such affiliation (e.g., a joint appointment or co-affiliated coauthors misattributed to one author) there is no way to tell which is “more recent” than another.

Field assignment works the same way but is weighted. OpenAlex tags each author with a set of topics, each carrying a `count` (roughly, the number of the author’s works classified under that topic) and a parent field. I sum `count` within each field and assign the author to whichever field carries the largest total weight: their modal field by publication volume, not simply their single most granular topic. The result is an `authors` table with one row per author: (`author_id`, `h_index`, `institution_id`, `institution_name`, `field`, `field_name`).

Two structural features of this procedure limit the impact of field classification errors. First, the paper operates at the *field* level (26 fields) rather than at the finer topic level ($\sim 4,500$ topics); OpenAlex’s topic classifier fine-tunes a multilingual BERT model on title, abstract, and citation data and achieves a top-1 accuracy of 0.53 (OpenAlex, 2024), but many topic-level errors stay within the same parent field, so field-level accuracy is plausibly higher, though this is an inference from the structure of the classification hierarchy rather than a directly measured field-level accuracy figure: no field-level validation of the classifier is presented here. Second, an author’s field is determined by the mode of their *entire publication record*, not by a single work: a researcher would need a substantial fraction of their career output systematically misclassified to a different field before their assignment would change.

2.3 The h-index and its successive generalization

Hirsch’s (2005) original index, here h_1 , is computed by OpenAlex per author from citation counts and taken as given. Schubert’s generalization (2007) lifts the same definition up a level of aggregation: a group of N members, each already carrying an h_{k-1} index, has an index $h_k = n$ if n of its members have $h_{k-1} \geq n$ and the remaining $N - n$ members have $h_{k-1} < n$:

Algorithm 1 Successive h-index of a group G , given a value $v(m)$ for each member $m \in G$

- 1: Sort G by v descending; let $r(m)$ be m ’s 1-indexed rank in that order
 - 2: **return** $\max\{r(m) : v(m) \geq r(m), m \in G\}$
-

I use this rank-based procedure, implemented as a window function in SQL ranking rows within each group and taking the largest rank whose value has not yet fallen below it, for every successive index in this paper: h_2 of an institution over the h_1 of its authors, and h_3 of a country over the h_2 of its institutions.

2.4 h_2 : institution and institution $\times$ field

h_2 is computed twice, over different groupings of the same author-level data. The field-specific table groups authors by (institution, field) and yields 288,300 pairs; institutions are matched on their numeric identifier only; institution and field *names* are resolved afterward via `arg_max`, so that spelling or formatting variants of the same name across snapshot partitions cannot silently split one institution into several groups. The institution-overall table ignores field entirely and applies the same procedure across all of an institution’s authors regardless of discipline, yielding one h_2 per institution (20,932 rows).

2.5 h_3 : country level

Institutions are mapped to countries using the `country_code` field of OpenAlex’s institutions snapshot (also read directly from S3, joined on institution ID). h_3 is then computed with the same rank-based procedure, once globally per country over the h_2 values of its institutions and once per (country, field) over the field-specific h_2 values, producing one file per field.

2.6 Efficiency

Mryglod, Holovatch, and Kenna (2022) showed that group-level h-indices grow with group size even when per-member quality is held constant. They were not studying successive h indices, but a similar argument makes raw h_2 and h_3 values unsuitable for comparing institutions or countries of unequal scale.

Egghe (2008) applied a Lotkaian framework to successive h indices. Suppose the article–citation frequency function follows Lotka’s law with exponent $\alpha_0 > 1$, and the author–article frequency function follows Lotka’s law with exponent $\alpha_1 > 1$. Egghe

showed that the distribution of h_1 values across authors is itself Lotkaian with compound exponent $\beta_1 = \alpha_0\alpha_1$. Applying the h-index formula for Lotkaian systems, the institution-level index satisfies

$$h_2 = S^{1/\beta_1} \quad (1)$$

where S is the number of authors at the institution. An analogous argument at the next level gives the distribution of h_2 values across institutions a Lotkaian exponent $\beta_2 = \alpha_0\alpha_1\alpha_2$, where α_2 is the Lotka exponent of the institution–author frequency function, so that

$$h_3 = R^{1/\beta_2} \quad (2)$$

where R is the number of institutions in the country. I define size-normalized efficiency metrics by dividing out these predicted scalings:

$$\varepsilon_2 = \frac{h_2}{S^{1/\beta_1}} \quad (3)$$

$$\varepsilon_3 = \frac{h_3}{R^{1/\beta_2}} \quad (4)$$

An institution with $\varepsilon_2 > 1$ outperforms the Lotkaian prediction for its headcount; one with $\varepsilon_2 < 1$ underperforms it.

The exponents β_1 and β_2 are estimated directly from the data rather than assembled from separately fit $\alpha_0, \alpha_1, \alpha_2$. I apply discrete maximum-likelihood tail estimation following [Clauset, Shalizi, and Newman \(2009\)](#), with $x_{\min}$ chosen to minimize the Kolmogorov–Smirnov distance between the empirical and fitted CDFs, to the distributions of h_1 across all authors and h_2 across all institutions, reading β_1 and β_2 as their fitted Lotka exponents. I then fit α_0, α_1 , and α_2 individually: α_0 from a stratified sample of per-article citation counts drawn from the OpenAlex works snapshot, α_1 from per-author work counts, and α_2 from per-institution author counts. Finally, I verify that $\alpha_0\alpha_1 \approx \beta_1$ and $\alpha_0\alpha_1\alpha_2 \approx \beta_2$ as a cross-check that the compounding-exponent structure of Egghe’s model holds in the OpenAlex dataset.

2.7 Supplementary metrics

Several secondary statistics are derived from the h_2 and h_3 tables to probe what the headline numbers obscure:

- **Gini coefficient**, computed on the within-field distribution of institutional h_2 , quantifies whether a field’s excellence is concentrated in one institution (Gini near 1) or broadly distributed (Gini near 0).
- **Breadth score** counts, per institution, how many of the 26 fields it ranks in the global top 10, top 50, and top 100.
- **Non-biomedical ranking** recomputes institution-level h_2 after excluding authors whose modal field is Medicine, Biochemistry/Genetics/Molecular Biology, Immunology and Microbiology, Neuroscience, or Health Professions, to test whether the overall ranking is, in substance, a ranking of medical school size.

2.8 Manual cross-validation

Because errors in OpenAlex’s underlying author disambiguation or citation counts would propagate directly into every successive index, I manually spot-check a targeted, non-random sample of authors rather than a random one, chosen at the points where an error would be both most likely and most consequential: authors sitting exactly at an institution’s h_2 threshold, where a ± 1 error in a single h_1 could change the institution’s rank; top authors in fields where OpenAlex’s topic classification is known to be thinner (e.g., Arts and Humanities); top authors at institutions with non-Anglophone naming conventions, where author-name disambiguation is more error-prone and where OpenAlex’s coverage is known to diverge unevenly from other databases by country, most severely for China (Zheng, Miao, Bu, & Larivière, 2025), including institutions suspected of affiliation-swapping for grant purposes (Bhattacharjee, 2011); and the small number of authors with the highest h_1 globally, as a sanity check at the extreme of the distribution. For each candidate, I resolve the author’s OpenAlex profile through the public API and compare it by hand against Web of Science, Scopus, and Google Scholar. This design is deliberately adversarial: it is built to surface errors, not to estimate how frequently they occur across the full population, so any outlier rate found in it bounds the qualitative failure mode (which authors, and by how much) rather than the population-wide error rate.

2.9 Sensitivity to disambiguation error

Because the cross-validation sample above cannot itself bound a population-wide error rate, I separately ask how much author-disambiguation error the paper’s conclusions could tolerate before changing using a Monte Carlo sensitivity analysis run directly on the full authors table. At a population-wide error rate p , each author is independently flagged with probability p ; a flagged author’s h_1 is multiplied by an inflation ratio drawn uniformly from the four ratios actually observed among the cross-validation outliers, simulating a merged-identity error of the same kind and magnitude found by hand. I recompute institution-level h_2 , country-level h_3 , and the field-specific h_2 rankings, and compare each to the unperturbed baseline via Spearman rank correlation and top-set overlap. Rather than picking a single rate to call a “worst case,” I run five independent trials at each of six rates, $p \in \{1\%, 5\%, 10\%, 20\%, 40\%, 80\%\}$, tracing out how ranking stability responds to the assumed amount of error from a plausible rate up to a deliberately implausible one, so that the paper’s tolerance for this kind of error is characterized by a curve rather than asserted at a single point.

2.10 Robustness to field assignment

The single-field assignment described in Subsection 2.2 asks “which field accounts for the largest share of an author’s publication volume?” and assigns each author exclusively to that field. To test whether field-specific h_2 rankings are stable under an alternative that allows multidisciplinary authors to count in multiple fields simultaneously, I re-run the field-specific h_2 computation after expanding each author’s field membership: at a threshold τ , an author is assigned to every field in which their OpenAlex topic-weight share is at least τ (where topic-weight is the same count-weighted

field sum used in the original assignment). I test two thresholds, $\tau \in \{20\%, 33\%\}$, spanning from liberal (an author need contribute only one fifth of their output to count in a field) to moderate (one third or more). For each threshold I recompute h_2 for every (institution, field) pair in the expanded author roster, using the same rank-based algorithm as the original, and compare to the baseline via Spearman rank correlation on shared (institution, field) pairs and by checking whether each field’s leading institution remains #1.

2.11 Robustness to institution assignment

The single-institution assignment described in Subsection 2.2 breaks ties among institutions that are equally “most recent” for a given author (identical latest affiliation year) by taking the numerically smallest OpenAlex institution identifier, silently discarding every other tied institution. To test how much this specific tiebreak choice, as distinct from current-affiliation attribution generally, is moving the institution- and country-level rankings, I identify every author with two or more educational institutions genuinely tied for their most recent affiliation year and re-assign that author to *every* tied institution, using the same inclusive-membership pattern as the multi-field robustness check in Subsection 2.10 (an author with k tied institutions is counted in full, not fractionally, at each of the k , rather than split $1/k$ per institution). This is not a literal fractional h-index; it is a sensitivity bound on how much of the ranking is attributable to the arbitrary tiebreak rule specifically, holding the underlying current-affiliation attribution method fixed. I recompute institution-level h_2 and country-level h_3 on the resulting inclusive multi-institution roster and compare each to the paper’s baseline single-pick assignment via Spearman rank correlation and top-set overlap, following the same procedure used for the field-robustness check above.

2.12 External validation against the Leiden Ranking

To assess whether ε_2 captures genuine research quality rather than a residual size artefact, I compare both raw h_2 and ε_2 against the CWTS Leiden Ranking Open Edition 2024 (van Eck, 2024), an OpenAlex-based institution-level bibliometric ranking based on mean normalised citation score (MNCS). Because it is built from the same underlying OpenAlex infrastructure as this paper, this comparison is a convergent check within a shared data source rather than an independent external validation; Subsection 2.13 below adds a genuinely independent comparison against a non-OpenAlex data source. I compute Spearman rank correlations $\rho(h_2, \text{MNCS})$ and $\rho(\varepsilon_2, \text{MNCS})$ on this matched set; p-values use the standard t -approximation $t = \rho\sqrt{(n-2)/(1-\rho^2)}$ evaluated against a t_{n-2} distribution.

2.13 External validation against Scimago Institutions Rankings

The Leiden comparison above shares OpenAlex source infrastructure with this paper’s own data, so it cannot rule out the possibility that both simply inherit the same coverage biases. As a genuinely independent check, I compare raw h_2 and ε_2 against the Scimago Institutions Rankings (SIR) 2026 “Overall Rank” export (SCImago Research Group, 2026), a Scopus-derived composite ranking (Research 50%, Innovation 30%,

Societal 20% sub-scores) built from an entirely different underlying citation database than either OpenAlex or the Web-of-Science-based Leiden Ranking. The public SIR export carries no ROR or OpenAlex identifier, so institutions cannot be joined on ID as in the Leiden comparison; instead I normalize institution names on both sides (lower-casing, stripping punctuation, and removing a small stopword list) and join on normalized name, excluding SIR-side names that collide with more than one SIR institution. This name-based join is necessarily less precise than the ID-based Leiden join, so I report match-rate diagnostics alongside the correlation result. I compute Spearman ρ between SIR’s Global Rank and both the raw h_2 rank and the ε_2 efficiency rank from the paper’s ≥ 100 -author efficiency table (Subsection 3.9), and repeat the comparison restricted to SIR’s “Universities” sector only, excluding the ministries, academies, hospitals, and companies that SIR also ranks, to check whether any mismatch is an artefact of comparing against a mixed-sector list.

2.14 Software and reproducibility

All analyses were implemented in Python 3.10 (NumPy, pandas, geopandas, matplotlib, duckdb, powerlaw, scipy). The data fetching, filtering, and processing are fully reproducible from the public code repository at <https://github.com/rupumped/successive-h-indices>.

3 Results

3.1 Manual cross-validation

Of the 30 authors sampled under the scheme described in Subsection 2.8, Table 1 summarizes agreement between OpenAlex’s `h_index` and the corresponding value I located by hand in Web of Science, Scopus, and Google Scholar. Coverage itself varies sharply: Scopus returned a value for 26 of the 30 authors, Google Scholar for 19, and Web of Science for 16.

Table 1 Agreement between OpenAlex and each external source, over authors for which that source returned a value. Diff is OpenAlex minus the external value.

Source	n	Mean diff	Mean abs. diff	Worst case
Web of Science	16/30	+24.5	25.7	Lei Jiang (+100)
Scopus	26/30	+8.0	12.4	Donald Small (+46)
Google Scholar	19/30	−22.7	30.7	A. Jafari (−85)

Scopus tracks OpenAlex most closely, both in coverage and in magnitude of disagreement. Google Scholar is systematically higher than OpenAlex, consistent with its broader crawl of preprints, theses, and other grey literature and its comparative permissiveness toward self-citation; this ordering (Google Scholar > Scopus > OpenAlex > Web of Science) reproduces cleanly among the six global top- h_1 sanity-check

authors, none of whom show a discrepancy exceeding 27% of their OpenAlex value in either direction.

Four authors, however, show disagreement well outside this range. I read this as a signature of author-disambiguation failure (multiple distinct people merged under one OpenAlex author ID) rather than of genuine cross-database disagreement:

- **Donald Small** (Johns Hopkins): $h_1^{\text{OpenAlex}} = 103$ against Scopus = 57 and Google Scholar = 59.
- **Karen S. Calhoun** (University of Michigan): $h_1^{\text{OpenAlex}} = 45$ against Scopus = 11 and Google Scholar = 13.
- **Lei Jiang** (University of Science and Technology of China): $h_1^{\text{OpenAlex}} = 215$ against Web of Science = 115, consistent with documented China-specific coverage discrepancies in OpenAlex (Zheng et al., 2025).
- **A. Jafari** (Sharif University of Technology): $h_1^{\text{OpenAlex}} = 139$, but Web of Science, Scopus, and Google Scholar report 41, 133, and 224 respectively.

All four fall in the non-Anglophone-institution or common-name strata of the sample, matching the concern raised in Subsection 2.8 that disambiguation is more error-prone for exactly these authors. Because such merges can only inflate an author’s apparent h_1 , and h_1 feeds directly into h_2 and h_3 , a handful of merged identities at the right institution could measurably distort an institution- or country-level ranking; I flag this as an open limitation rather than one this sample can size definitively. This sample is deliberately adversarial, chosen to find errors, not to estimate how often they occur. The 4/30 (13%) outlier fraction should not be read as an estimate of how much of the full dataset is affected. The same pattern, disambiguation precision concentrated lower for common names and non-Anglophone institutions, appears in an independent benchmark of 2,700 scholars sampled systematically across three regions (Zhao & Chen, 2025), corroborating that this is a geographically and onomastically concentrated failure mode rather than an artifact of this particular sample. OpenAlex’s maintainers report an ongoing correction effort targeting exactly this kind of over-merging, with corresponding-author identification precision rising from 0.60 to 0.92 (F1 0.72 to 0.90) after a 2026 rework of the underlying author-matching model (Demes, 2026). Subsection 3.12 quantifies how much a population-wide error rate at or below the adversarial sample’s raw rate would actually move this paper’s rankings.

The four Arts and Humanities threshold authors sampled (Gillian Ramchand and G. Herbert Fowler at Oxford; D.H. Mellor and James Raven at Cambridge) illustrate the range of under-coverage across major citation databases: Ramchand and Fowler returned no value in any of the three external sources, while Mellor (OA $h_1 = 28$) and Raven (OA $h_1 = 24$) appeared in Scopus at roughly half their OpenAlex values ($h_1 = 13$ and 12, respectively); for Raven, Google Scholar (26) brackets the OpenAlex value closely. This is consistent with the well-documented weaker coverage of book-based humanities scholarship in all major citation databases (Nederhof, 2006) and suggests that h-index comparisons in these fields should be read with correspondingly less confidence than in the sciences.

3.2 The overall h_2 ranking

Table 2 lists the ten institutions with the highest institution-level h_2 , computed over all 20,932 institutions and 30.0 million authors described in Section 2.

Table 2 Top 10 institutions by overall h_2 .

Institution	h_2	Authors
Harvard University	137	110,979
Stanford University	107	65,197
University of Cambridge	105	52,307
Cornell University	104	70,718
Johns Hopkins University	103	80,028
University of Oxford	103	57,110
University of Washington	102	71,270
University of North Carolina at Chapel Hill	101	64,478
University College London	100	77,482
University of California, San Francisco	100	50,351

This broadly tracks the US News and World Report ranking, where Harvard also places 1st, Stanford 3rd, and Cambridge 5th. As I show below, however, the overall ranking is a considerably less interesting object than it first appears.

3.3 Field-level heterogeneity

Splitting h_2 by (institution, field) reveals that the overall ranking is deceptive: several institutions that never come close to the top of Table 2 rank first globally in their field by this measure. Table 3 lists a sample of these cases.

University of Antwerp is the counterexample that proves the rule holds both ways: it is #14 overall ($h_2 = 98$) and #1 in Physics and Astronomy ($h_2 = 96$), a broad powerhouse that also ranks first in a specific, non-biomedical field. Universidade Estadual de Campinas (UNICAMP) places 6th in Dentistry behind São Paulo, placing Brazil among the top six globally in Dentistry by this measure despite the country appearing nowhere near the top of any general-purpose ranking (Vaccaro, Sánchez-Núñez, & Witt-Rodríguez, 2022). China’s cluster is not one specialty but four related applied-science fields: University of Science and Technology of China (USTC) ranks first in Materials Science, Engineering, and Energy by field h_2 , while its sister institution, the University of Chinese Academy of Sciences, ranks first in Chemical Engineering.

Not every field supports this kind of comparison equally well. Decision Sciences is the thinnest field in the dataset: Stanford’s global #1 is $h_2 = 22$, built on only 285 Stanford authors against 140,147 authors field-wide, versus Medicine’s $h_2 = 119$ built on 7.58 million authors field-wide. h_2 is not meaningfully comparable across fields of such different size, only within a field, and Decision Sciences in particular may be too sparse an OpenAlex topic category for its ranking to carry much signal.

Table 3 Selected field leaders that rank well outside the overall top 10. “Overall” columns give the institution’s rank and h_2 in Table 2’s full ranking; “Field h_2 ” is the institution’s h_2 restricted to that field.

Field	Leader	Ov. rank	Ov. h_2	Field h_2
Agricultural & Biological Sciences	Wageningen	172	71	56
Environmental Science	Wageningen	172	71	56
Pharmacology, Toxicology & Pharmaceutics	Shenyang Pharmaceutical University	535	53	24
Computer Science	Carnegie Mellon University	175	70	56
Materials Science	USTC	39	89	59
Engineering	USTC	39	89	62
Energy	USTC	39	89	52
Chemical Engineering	Univ. of Chinese Academy of Sciences	91	80	23
Dentistry	Universidade de São Paulo	235	66	34

3.4 Robustness of field-specific rankings to multi-field assignment

The single-field assignment described in Subsection 2.2 asks which field accounts for the plurality of an author’s publication volume, and assigns them exclusively there. To test how much this choice moves the field-level rankings, I re-run the field-specific h_2 computation after allowing each author to count in every field for which their OpenAlex topic-weight share is at least τ in accordance with Subsection 2.10.

At $\tau = 33\%$ (an author must contribute at least one third of their output to count in a field), the (institution, field) pair count rises by 9.3% (333,132 $\rightarrow$ 364,032 pairs). The Spearman rank correlation between original and multi-field h_2 values, computed on the inner join of shared pairs, is $\rho = 0.963$: field rankings are highly stable. Of the 27 fields’ leading institutions (26 fields, but Social Sciences has two institutions tied for the top h_2 value, yielding 27 leader entries), 23 are unchanged. The Wageningen Ag&Bio leadership (Table 3) survives and Carnegie Mellon’s Computer Science lead survives. The one notable exception is Environmental Science: the University of Colorado Boulder overtakes Wageningen at this threshold, because Boulder has a larger body of researchers who do substantial but not modal Environmental Science work. The two methods answer different questions. The single-field assignment identifies institutions where Environmental Science is researchers’ *primary* field; the multi-field

assignment identifies institutions with the most researchers for whom it is a *substantial* field.

At $\tau = 20\%$ (a more aggressive threshold admitting all authors with one fifth or more of their output in a field), the pair count rises by 20.7% and $\rho = 0.907$; 18 of 27 field leaders survive. Carnegie Mellon’s Computer Science lead also falls at this threshold, with Stanford leading instead, consistent with the same mechanism of Stanford having a larger population of researchers with at least 20% CS output, even if CS is not their modal field. The Wageningen Ag&Bio and Environmental Science results are unchanged from the 33% case. These findings confirm that the overall field ranking structure is broadly stable under multi-field assignment but that the field-leadership claims for smaller, highly specialised institutions depend on how narrowly “counts in this field” is defined. Given this sensitivity, the field-leadership claims in Table 3 should be read as reflecting depth of specialisation under the paper’s single-field assignment.

3.5 Robustness of institutional rankings to affiliation-tie assignment

The single-institution assignment described in Subsection 2.2 breaks ties among an author’s equally-most-recent educational institutions by taking the smallest OpenAlex institution identifier. Roughly 3,747,235 of the 29,970,900 qualifying authors (12.5%) have a genuine tie: two or more educational institutions with an identical latest affiliation year. Re-assigning each such author to every tied institution, following the inclusive multi-membership procedure of Subsection 2.11, expands the (author, institution) roster from 29,970,900 pairs to 35,291,037 (+17.8%) and adds 271 institutions to the 20,932 that clear the data filters.

Institution-level h_2 is highly stable under this re-assignment: on the 20,932 institutions shared between the baseline and multi-institution rosters, the Spearman rank correlation is $\rho = 0.9658$, with top-10 overlap of 80% and top-20 overlap of 70%. Country-level h_3 is even more stable, as expected since institution→country mapping is a fixed lookup untouched by this re-assignment: $\rho = 0.9965$ across all 214 shared countries, with top-10 overlap of 100% and top-20 overlap of 95%.

Saudi Arabia’s efficiency figure, flagged elsewhere in this paper (Subsection 3.11) as a plausible artefact of funding-driven affiliation listing (Bhattacharjee, 2011), moves under this check, but in the wrong direction to be explained by the tiebreak heuristic: ε_3 rises from 6.692 to 7.495 once tied authors are counted at every tied institution, rather than falling as would be expected if the arbitrary tiebreak were artificially inflating Saudi institutions’ apparent headcount advantage. This rules out the current-affiliation tiebreak mechanism specifically as the source of Saudi Arabia’s efficiency anomaly; it does not identify what the true mechanism is, and is consistent with the funding-affiliation explanation already offered elsewhere in the paper, which this check cannot test directly.

Taken together, these results indicate that the arbitrary tiebreak rule used for single-institution assignment is a comparatively minor contributor to ranking instability: it moves institution rankings by a bounded, measurable amount and country rankings by very little. This does not, however, resolve the broader concern that a

researcher’s entire career h-index is credited to wherever they are currently affiliated, tied or not; that concern is addressed in this paper by narrowing its claims rather than by further re-attribution (Subsection 2.2 and the Limitations paragraph of the Conclusion).

3.6 Concentration within fields

Table 4 reports, for every field, the Gini coefficient of institutional h_2 defined in Section 2.7, together with the field’s institution count, leading (maximum) h_2 , and mean h_2 .

Table 4 Gini coefficient of institutional h_2 within each field, sorted from most to least concentrated.

Field	Institutions	Max h_2	Mean h_2	Gini
Physics and Astronomy	12,739	96	6.11	0.605
Biochemistry, Genetics & Molecular Biology	15,404	104	6.61	0.578
Neuroscience	11,646	64	4.41	0.574
Immunology and Microbiology	9,838	67	3.97	0.561
Materials Science	13,097	59	6.36	0.560
Earth and Planetary Sciences	12,447	60	4.30	0.557
Chemistry	12,103	41	5.02	0.541
Environmental Science	16,603	56	5.54	0.541
Engineering	17,742	62	6.91	0.533
Medicine	19,119	119	8.58	0.528
Agricultural and Biological Sciences	16,391	56	5.61	0.522
Energy	8,971	52	3.38	0.515
Psychology	15,745	58	3.64	0.513
Computer Science	16,736	56	5.08	0.512
Dentistry	8,303	34	3.01	0.500
Mathematics	11,092	36	3.42	0.491
Economics, Econometrics and Finance	13,923	55	3.30	0.482
Business, Management and Accounting	14,623	32	3.59	0.481
Health Professions	14,840	41	2.88	0.465
Social Sciences	19,242	48	4.57	0.461
Pharmacology, Toxicology and Pharmaceutics	7,219	24	2.61	0.459
Chemical Engineering	5,874	23	2.48	0.446
Arts and Humanities	15,390	30	2.78	0.443
Nursing	8,958	31	2.37	0.433
Decision Sciences	10,409	22	2.37	0.423
Veterinary	4,678	28	2.04	0.422

Physics and Astronomy is the most unequal field: University of Antwerp’s $h_2 = 96$ is 15.7 times the field mean of 6.1. Veterinary and Decision Sciences are the most evenly distributed (Gini ≈ 0.42). Medicine sits in the middle of the table (0.528) despite

Harvard’s large absolute lead, because Medicine, unlike Physics and Astronomy, has thousands of institutions with meaningful h_2 rather than a single dominant outlier.

3.7 The overall ranking as an implicit biomedical ranking

Medicine alone accounts for 7.58 million of the dataset’s 30.0 million authors (25%), and, together with Biochemistry/Genetics/Molecular Biology, Immunology and Microbiology, Neuroscience, and Health Professions, the five core biomedical fields dwarf every other field’s author pool. Table 5 recomputes institution-level h_2 after excluding authors in these five fields entirely to test whether the overall ranking in Table 2 is, in substance, a ranking of medical school size.

Table 5 Top 10 institutions by non-biomedical h_2 , alongside their position in the unrestricted overall ranking. Δ rank is overall rank minus non-biomedical rank; positive means the institution moves up once biomedical fields are excluded.

Institution	Rank	h_2	Overall rank	Overall h_2	Δ rank
Harvard University	1	100	1	137	0
University of Antwerp	2	97	14	98	+12
California Institute of Technology	3	94	15	97	+12
Stanford University	4	92	2	107	−2
Université Paris-Saclay	5	91	22	95	+17
Massachusetts Institute of Technology	6	89	25	94	+19
University of Geneva	7	88	29	93	+22
Univ. of Science and Technology of China	8	88	39	89	+31
ETH Zurich	9	87	36	89	+27
University of Cambridge	10	87	3	105	−7

Harvard remains #1 even with its biomedical authors removed, confirming the claim in Subsection 3.3 that it is genuinely excellent rather than merely large. But the rest of the table reshuffles substantially: Antwerp, Caltech, Paris-Saclay, MIT, Geneva, USTC, and ETH Zurich all gain 12 to 31 places once Medicine and its four sibling fields are set aside, while Stanford and Cambridge, both anchored by large medical schools, fall by 2 and 7 places respectively. The University of Chinese Academy of Sciences shows the same pattern one tier down: 91st overall but 23rd non-biomedical. Because medical fields have by far the largest, most-cited author pools, they structurally dominate any university-wide aggregate; an institution elite in non-medical fields, such as USTC or Carnegie Mellon, can never reach the top of the overall list no matter how good it is, simply because it lacks a large medical school. The field-level and non-biomedical breakdowns matter more than the unrestricted university-level number.

3.8 Breadth versus depth

Table 3 showed institutions that lead a single field while sitting far outside the overall top 100. Table 6 shows the opposite pattern: institutions whose strength is broad

rather than deep, measured by how many of the 26 fields they place in the global top 50.

Table 6 The nine institutions with the most top-50 field placements out of 26 fields.

Institution	Top 50	Top 10	Overall rank	Overall h_2
University of California, Los Angeles	20	8	13	99
University of Michigan	20	7	26	94
Harvard University	19	13	1	137
University of Oxford	19	8	6	103
University of Washington	17	8	7	102
University of Toronto	17	7	11	100
Cornell University	17	6	4	104
Univ. of North Carolina at Chapel Hill	17	5	8	101
University of Cambridge	17	4	3	105

UCLA and the University of Michigan share the greatest breadth by this measure, each placing in the global top 50 in 20 of 26 fields; UCLA leads 8 outright while Michigan leads 7. Harvard still leads more fields outright (13 of 26) than any other institution, but four of the nine broadest institutions by this top-50 count, Toronto, UNC Chapel Hill, Washington, and Cambridge, place fewer than half their 17 top-50 fields inside the top 10, meaning most of their breadth comes from being solidly good across many fields rather than dominant in any one. At the other extreme, Boston University (rank 23, $h_2 = 94$) achieves its overall strength with zero top-10 field placements, through scale and consistency across many merely-good departments rather than fielding the single best group in any one discipline.

3.9 Institution efficiency

Model validation

Figure 1 plots institution-level h_2 against the Egghe-predicted value S^{1/β_1} for all institutions with at least 10 OpenAlex authors. A no-intercept regression, appropriate because $N = 0$ authors forces $h_2 = 0$, gives slope $c = 1.945$ and uncentered $R^2 = 0.911$: S^{1/β_1} alone explains 91% of the variance in h_2 , confirming that Egghe’s power-law functional form describes the data. The efficiency metric $\varepsilon_2 = h_2/S^{1/\beta_1}$ captures each institution’s deviation from this overall trend.

The cross-check on Egghe’s compounding structure fails sharply, however. Egghe’s derivation requires that the Lotka tail exponent of the h_1 distribution equals the product $\alpha_0\alpha_1$; these separate fits give $\alpha_0 = 2.742$ and $\alpha_1 = 2.996$, whose product is 8.21, against the directly-measured $\beta_1 = 2.963$: a discrepancy of 177%. The analogous prediction at the next level ($\beta_2 = \beta_1\alpha_2$) is similarly violated: $2.963 \times 1.994 = 5.909$ against the directly-measured $\beta_2 = 2.998$ (97% discrepancy).

The failure of the compounding formula does not, however, undermine the efficiency metric. The two claims play different roles. The compounding formula ($\beta_1 =$




Fig. 1 Institution-level h_2 versus S^{1/β_1} (Egghe’s predicted scaling, $\beta_1 = 2.963$). No-intercept regression: $h_2 = 1.945 S^{1/\beta_1}$, $R^2 = 0.911$ (uncentered).

$\alpha_0\alpha_1$) is a *theoretical prediction* about the relationship between separately-fitted exponents; it is tested here and found wanting. The power-law scaling $h_2 \propto S^{1/\beta_1}$ is an *empirical finding* ($R^2 = 0.911$) that does not depend on the compounding formula being correct. β_1 is estimated directly from the distribution of h_1 values across authors, not derived from $\alpha_0\alpha_1$, so the 177% discrepancy in the compounding formula has no bearing on the exponent actually used in ε_2 . This is checked against the Leiden Ranking and, **more independently, against Scimago Institutions Rankings**, later in this subsection, **though as discussed there, raw h_2 tracks both external rankings more closely than ε_2 does, so those comparisons should be read as description of the efficiency metric’s behavior rather than as confirmation of it.**

Efficiency rankings

The efficiency metric $\varepsilon_2 = h_2/S^{1/\beta_1}$ addresses Mryglod et al.’s (2022) concern directly: because S^{1/β_1} accounts for 91% of the variance in h_2 , dividing it out isolates the residual variation that is not explained by headcount alone.

Table 7 lists the top 10 institutions by ε_2 among the 12,026 with at least 100 authors. The top of the ranking is dominated by small and highly specialized institutions. Focused research centers and specialized colleges recur throughout the top 20.

Four of the ten institutions in Table 7 — Bellevue University, Imperial Valley College, City College of San Francisco, and Frederick Community College — are primarily-teaching institutions whose reported h_2 is implausible on its face for their type. Manually inspecting each institution’s highest- h_1 contributing authors finds the mechanism: rather than genuine small research populations, these authors each list

Table 7 Top 10 institutions by efficiency $\varepsilon_2 = h_2/S^{1/\beta_1}$ ($\beta_1 = 2.963$) among those with at least 100 authors ($n = 12,026$).

Institution	ε_2	h_2	Authors	h_2 rank
Bellevue University	6.24	54	590	497
Augustana University	5.61	64	1,332	264
Center for Behavioral Brain Sciences	5.48	31	168	1,908
Barcelona Inst. of Science and Technology	5.43	53	839	520
Imperial Valley College	5.33	62	1,412	319
Saint Camillus International Univ. of Health Sci.	5.30	39	364	1,228
City College of San Francisco	4.92	43	605	934
Frederick Community College	4.81	37	416	1,331
Académie Nationale de Médecine	4.72	39	512	1,192
Natl. Interuniversity Consortium of Materials Sci.	4.67	40	571	1,150

an implausible number of distinct lifetime educational institutions (46–340, against a median of 4–7 across the rest of each institution’s roster), spanning unrelated countries, disciplines, and career eras within a single record — the signature of OpenAlex author-disambiguation merge failure (several distinct researchers collapsed into one author identifier) rather than authors genuinely affiliated with these institutions. Excluding authors with 20 or more distinct lifetime educational institutions and recomputing h_2 for these four institutions in isolation gives 28 (Bellevue, from 54), 55 (Imperial Valley, from 62), and 35 (City College of San Francisco, from 43); Frederick Community College’s h_2 is unchanged at 37, since only one of its 416 assigned authors meets this threshold and removing it does not cross the rank-37 value. This substantially closes the gap for three of the four institutions, though not all the way to a value a two-year teaching institution’s research output would plausibly support on its own, indicating the merge-artifact pattern is not fully captured by this single threshold. I do not apply this exclusion to the paper’s other tables: a systematic evaluation of how many institutions across the full dataset carry the same pattern is beyond the scope of the present revision, and the top of Table 7 should be read with this caveat rather than as a validated efficiency finding.

Table 8 restricts to the 1,339 institutions with at least 5,000 authors. The leading institutions are now a mix of specialized research institutes like Rockefeller University ($\varepsilon_2 = 3.77$, biomedical), the Institute of Cancer Research ($\varepsilon_2 = 3.57$), and the London School of Hygiene & Tropical Medicine ($\varepsilon_2 = 3.44$) and comprehensive universities like the University of Antwerp ($\varepsilon_2 = 3.72$, $h_2 = 98$, efficiency rank 112). The University of California System, treated here as a single institution under its OpenAlex identifier, is the only entry in the top 100 on both metrics simultaneously (h_2 rank 73, efficiency rank 67). Among the comprehensive elite research universities that dominate the raw h_2 table, Caltech (efficiency rank 367), the University of Geneva (344), and Université Paris-Saclay (465) are the most efficient, each achieving $h_2 \geq 93$ with under 25,000 authors.

Table 8 Top 10 institutions by efficiency $\varepsilon_2 = h_2/S^{1/\beta_1}$ ($\beta_1 = 2.963$) among those with at least 5,000 authors ($n = 1,339$).

Institution	ε_2	h_2	Authors	h_2 rank
University of California System	3.93	83	8,242	73
Rockefeller University	3.77	70	5,637	179
University of Antwerp	3.72	98	15,759	14
Institute of Cancer Research	3.57	66	5,570	230
University of San Diego	3.56	64	5,097	286
California University of Pennsylvania	3.49	76	9,011	121
University of Wuppertal	3.48	68	6,515	209
London School of Hygiene & Tropical Medicine	3.44	65	5,903	249
Royal Holloway, University of London	3.39	61	5,102	339
Vita-Salute San Raffaele University	3.37	71	8,171	171

External validation against the Leiden Ranking

Matching the 1,506 Leiden Ranking Open Edition 2024 universities to OpenAlex’s dataset yields 1,495 institutions with both an MNCS value and an institution-wide h_2 from OpenAlex. The Spearman rank correlation between raw h_2 and MNCS is $\rho = 0.634$ ($p \approx 6 \times 10^{-170}$); between ε_2 and MNCS it is $\rho = 0.580$ ($p \approx 5 \times 10^{-135}$). **Because the Leiden Open Edition is itself built on OpenAlex, this agreement is evidence of convergence within a shared data infrastructure rather than independent validation; Subsubsection 3.9 below repeats the comparison against a source with no OpenAlex dependency.**

External validation against Scimago Institutions Rankings

Joining the Scimago Institutions Rankings (SIR) 2026 Overall Rank export (10,827 institutions, Scopus-derived) to the paper’s ≥ 100 -author efficiency table (12,026 institutions) on normalized institution name, since the SIR export carries no ROR or OpenAlex identifier, yields 3,659 matched institutions: 33.8% of SIR’s list and 30.4% of the paper’s efficiency table. This is a materially less precise join than the ID-based Leiden match above, so it is read alongside that caveat.

On the matched set, the Spearman rank correlation between raw h_2 and SIR’s Global Rank is $\rho = 0.752$ ($p \approx 0$, $n = 3,659$); between ε_2 and SIR’s Global Rank it is $\rho = 0.338$ ($p \approx 2 \times 10^{-98}$). Restricting the match to SIR’s “Universities” sector only (5,491 of SIR’s 10,827 institutions, excluding ministries, academies, hospitals, and companies that SIR also ranks) gives essentially the same result ($\rho = 0.755$ raw, $\rho = 0.332$ efficiency), so the gap between the two correlations is not an artefact of comparing an academic h_2 ranking against a mixed-sector SIR list.

This is the same pattern as the Leiden comparison above, now against a source with no shared data infrastructure with OpenAlex at all: raw h_2 tracks the external ranking (Leiden’s MNCS, SIR’s Global Rank) more closely than the size-normalized efficiency metric ε_2 does. I report both correlations as description rather than as validation of the efficiency normalization; a plausible mechanistic reading is that external

composite rankings like SIR’s Global Rank and Leiden’s MNCS are themselves influenced by scale-correlated factors (institutional prominence, resource concentration) that raw h_2 partially captures and that dividing out size by construction removes. That both external comparisons show this same direction, using different data sources and different underlying constructs, argues against either result being a coincidence of one particular external ranking.

3.10 h_3 : country level

Aggregating institutional h_2 up to the country level via the same rank-based procedure yields h_3 . Table 9 lists the ten highest- h_3 countries, and Figure 2 maps h_3 for all 214 countries with at least one qualifying institution.

Table 9 Top 10 countries by h_3 .

Country	h_3	Institutions
United States	70	4,242
China	55	1,078
Germany	49	505
United Kingdom	47	548
Japan	46	1,358
Italy	45	229
France	41	394
Spain	39	156
South Korea	37	358
Australia	36	120

The United States leads by a wide margin, but Italy’s 45 on only 229 institutions and Japan’s 46 on 1,358 are a first hint that raw h_3 rewards national research-system size as well as depth, a point to which I return in Subsection 3.11.

3.11 Country efficiency

Figure 3 repeats the size-scaling analysis at the country level, plotting h_3 against R^{1/β_2} for each country. The no-intercept fit gives slope $c = 4.064$ and uncentered $R^2 = 0.922$: institution count alone explains 92% of the variance in h_3 across countries.

Table 10 ranks the 151 countries with at least 5 qualifying institutions by efficiency $\varepsilon_3 = h_3/R^{1/\beta_2}$. The Italy–Japan contrast foreshadowed in Subsection 3.10 is now quantified: Italy ($\varepsilon_3 = 7.35$, rank 1) achieves $h_3 = 45$ from only 229 institutions, while Japan ($\varepsilon_3 = 4.15$, rank 44) reaches $h_3 = 46$ from 1,358, virtually the same headline number but at 6× the institutional scale. Australia (rank 2, $\varepsilon_3 = 7.29$, 120 institutions) and Spain (rank 3, $\varepsilon_3 = 7.24$, 156 institutions) round out a top tier of countries with compact research systems that place well on this metric. Germany (rank 6), the United Kingdom (rank 7), and France (rank 8) are the only countries that place in the top 10 on both raw h_3 and efficiency; placing well on both measures simultaneously is a weaker claim than confirming that the underlying research depth is genuine, since




Fig. 2 h_3 by country. Countries with no qualifying institution are shown in grey.

neither raw h_3 nor ε_3 is constructed to rule out the kind of affiliation-attribution distortion discussed below for Saudi Arabia. Saudi Arabia is a notable outlier at 4th by efficiency ($h_3=25$ on just 52 institutions) despite ranking only 20th by raw h_3 ; as flagged in Subsection 2.8, this is plausibly an artifact of high-h-index researchers listing Saudi institutions as secondary affiliations in exchange for funding (Bhattacharjee, 2011) rather than a genuine concentration of resident research talent. Saudi Arabia is simply the case where this risk is most visible in the numbers; the multi-institution robustness check in Subsection 3.5 rules out the current-affiliation tiebreak heuristic specifically as its cause (Saudi Arabia’s ε_3 rises, not falls, under that check) but does not establish that other countries’ efficiency ranks are free of comparable affiliation-attribution risk, including the possibility of secondary or funding-driven affiliation listing, which this paper cannot test directly.

The United States falls to efficiency rank 34 ($\varepsilon_3 = 4.32$) despite its raw lead ($h_3 = 70$, rank 1). With 4,242 qualifying institutions, nearly $8\times$ more than China and $18\times$ more than Italy, even a large absolute h_3 advantage is absorbed by the size denominator. China (rank 10, $\varepsilon_3 = 5.36$, 1,078 institutions) is more compact and appears above the US on this metric.

3.12 Sensitivity to disambiguation error

Table 11 reports, for six simulated population-wide error rates spanning $p = 1\%$ to a deliberately implausible $p = 80\%$ (Subsection 2.9), the Spearman rank correlation and top-set overlap between the perturbed and unperturbed rankings, averaged over five trials, together with how often the paper’s two headline field-leadership claims survive




Fig. 3 Country-level h_3 versus R^{1/β_2} (Egghe’s predicted scaling, $\beta_2 = 2.998$). No-intercept regression: $h_3 = 4.064 R^{1/\beta_2}$, $R^2 = 0.922$ (uncentered).

Table 10 Top 10 countries by efficiency

$\varepsilon_3 = h_3/R^{1/\beta_2}$ ($\beta_2 = 2.998$) among those with at least 5 institutions ($n = 151$).

Country	ε_3	h_3	Institutions	h_3 rank
Italy	7.35	45	229	6
Australia	7.29	36	120	10
Spain	7.24	39	156	8
Saudi Arabia	6.69	25	52	20
Sweden	6.45	25	58	21
Germany	6.14	49	505	3
United Kingdom	5.74	47	548	4
France	5.59	41	394	7
Greece	5.36	21	60	31
China	5.36	55	1,078	2

the perturbation intact. The cross-validation sample’s own raw outlier fraction ($4/30 \approx 13\%$ across both waves, Subsection 3.1) is included as an extra row for reference, though it carries no special status in this analysis.

Institution- and country-level rank correlations decline only modestly with p and then plateau rather than continuing toward zero: Spearman $\rho(h_2)$ falls from 0.999 at $p = 1\%$ to 0.991 at $p = 20\%$ and then stays flat at 0.988 all the way out to $p = 80\%$; $\rho(h_3)$ shows the same pattern, from 0.9997 down to 0.994. This follows from how h_2 and h_3 are computed: each is an order statistic (the largest rank whose value has not yet fallen below it), so what matters for a given institution or country is whether a

noisy author sits near *its own* threshold, not how many noisy authors exist elsewhere in the dataset; and because the perturbation is bounded in magnitude (at most $\approx 4\times$, the largest of the four observed ratios) and applied independently and identically everywhere, raising p further increasingly inflates competing institutions’ threshold authors by similar amounts, which cancels out much of the relative shift.

The field-specific leadership claims for Wageningen (Agricultural & Biological Sciences, Environmental Science) and Carnegie Mellon (Computer Science) are noisier with five trials per rate: each survives in 3 to 5 of 5 trials at every rate tested. Wageningen’s Environmental Science and Carnegie Mellon’s Computer Science leads are each knocked out in exactly two trials across the entire grid; Wageningen’s Agricultural & Biological Sciences lead is the most sensitive, failing in four trials (one at $p = 10\%$, one at $p = 13\%$, and two at $p = 20\%$). Italy and Japan’s near-parity in h_3 despite a $6\times$ difference in institutional scale is preserved throughout the tested range: both values rise together under added noise (Italy 45 at $p = 1\%$ to 66–67 at $p = 80\%$; Japan 46 to 86–88 over the same range), and while their ratio drifts from about 0.96 down to about 0.76 as p increases, even at that extreme it remains more than four times tighter than the underlying 1:6 gap in institutional count.

Reading these together: the paper’s overall institution and country rankings, and its country-comparison claims, are robust to author-disambiguation error at every rate tested, including rates far beyond what any plausible reading of the cross-validation sample would suggest. The field-leadership claims are more sensitive, as single-institution results necessarily are, and should be read with correspondingly more caution than the paper’s other claims.

Table 11 Sensitivity of h_2 and h_3 rankings to simulated disambiguation error, by population-wide error rate p (5 trials per rate). Overlap is computed on the top 20 institutions (by h_2) and top 10 countries (by h_3). The last three columns give the fraction of trials in which the named institution keeps rank #1 in the stated field.

p	$\rho(h_2)$	Top-20	$\rho(h_3)$	Top-10	Wag. Ag&Bio #1	Wag. EnvSci #1	CMU CS #1
1%	0.999	93%	0.9997	94%	5/5	5/5	5/5
5%	0.997	84%	0.999	96%	5/5	5/5	5/5
10%	0.994	78%	0.998	90%	4/5	5/5	5/5
13%*	0.993	78%	0.998	90%	4/5	5/5	5/5
20%	0.991	80%	0.997	90%	3/5	4/5	3/5
40%	0.988	79%	0.996	90%	5/5	4/5	5/5
80%	0.988	83%	0.994	90%	5/5	5/5	5/5

*cross-validation sample’s raw outlier fraction, included for reference only.

4 Conclusion

Schubert closed his 2007 note wondering whether h_2 of institutes would prove more interesting than h_2 of journals, absent the researcher–institute–country data needed to check. Nineteen years and 30.0 million OpenAlex authors later, the answer is yes, but not in the way a single top-10 table would suggest.

The overall h_2 ranking is the least interesting object in this paper. Because Medicine alone accounts for 25% of the author pool, the unrestricted university-wide aggregate is, in substance, a ranking of medical school size. The field decomposition reveals a completely different map of global research strength: Wageningen ranks first in Agricultural & Biological Sciences and Environmental Science by field h_2 while ranking 172nd overall, though the Environmental Science lead holds only under the paper’s single-field assignment and not under multi-field assignment (Subsection 3.4); Carnegie Mellon ranks first in Computer Science by field h_2 from 175th, **though, like the Environmental Science lead, this too holds only under the paper’s single-field assignment: Carnegie Mellon’s Computer Science lead falls to Stanford under the more aggressive 20% multi-field threshold (Subsection 3.4)**; the University of Science and Technology of China ranks first in Materials Science, Engineering, and Energy from 39th; and Shenyang Pharmaceutical University ranks first in Pharmacology from 535th. Brazil offers perhaps the sharpest illustration of what the overall ranking hides: the Universidade de São Paulo and the Universidade Estadual de Campinas rank among the top six globally in Dentistry by field h_2 despite neither country nor either institution appearing anywhere near the top of any general-purpose ranking. The field-level h_2 is the unit that should be the default for comparison, not the aggregate.

On the methodological side, Egghe’s power-law functional form is empirically well-supported: **author count alone explains 91% of the variance in h_2 , and institution count alone explains 92% of the variance in h_3** . However, the theoretical compounding prediction $\alpha_0\alpha_1 \approx \beta_1$ fails by 177%. The model is empirically useful; the underlying theory needs a different reading. **Two external comparisons bear on whether the efficiency normalisation retains a genuine research-quality signal. Against the CWTS Leiden Ranking Open Edition 2024 (van Eck, 2024), a ranking that shares OpenAlex’s underlying data infrastructure and so functions as a convergent check rather than an independent validation, both h_2 and ε_2 correlate with the Leiden mean normalised citation score across 1,495 matched institutions, but raw h_2 ($\rho = 0.634$) correlates more strongly than ε_2 ($\rho = 0.580$). Against Scimago Institutions Rankings 2026 (SCImago Research Group, 2026), a Scopus-derived ranking independent of OpenAlex, the same pattern recurs and is more pronounced: raw h_2 correlates at $\rho = 0.752$ against SIR’s Global Rank versus $\rho = 0.338$ for ε_2 , on a smaller, name-matched set of 3,659 institutions (Subsections 2.12–2.13, Subsubsections 3.9–3.9). Both comparisons show raw, scale-dominated h_2 tracking external rankings more closely than the size-normalised ε_2 does, plausibly because those external composites are themselves influenced by scale-correlated factors that dividing out size removes; this is reported as a description of the metrics’ behaviour, not as validation of the efficiency normalisation’s superiority.**

Several limitations constrain interpretation. Author disambiguation is the most consequential open problem: the four outlier cases identified in cross-validation (notably Lei Jiang, $h_1^{\text{OpenAlex}} = 215$ against Web of Science = 115) confirm that merger errors cluster in non-Anglophone and common-name strata, exactly the authors whose inflated h_1 values would propagate most damagingly into institutional and country-level rankings. That cross-validation sample was deliberately adversarial, however, and cannot by itself say how much of the full **30.0-million-author** population is affected; a

Monte Carlo sensitivity analysis run at population-wide error rates from 1% to a deliberately implausible 80% (Subsections 2.9, 3.12) finds the risk bounded rather than fatal, with institution- and country-level rank correlations declining only modestly before plateauing ($\rho \geq 0.988$) rather than degrading toward noise at any rate tested, and the paper’s country-comparison claims holding throughout. The field-leadership claims are noisier and should be read with correspondingly more caution than the paper’s other claims. Arts and Humanities are systematically under-covered in all major citation databases: both threshold authors sampled at Oxford returned no value in Web of Science, Scopus, or Google Scholar, and h-index comparisons in book-based fields should be read with correspondingly less confidence. Field-specific rankings are also sensitive to how field membership is defined: a multi-field assignment (authors counted in every field comprising $\geq 33\%$ of their output) yields highly stable overall rankings ($\rho = 0.963$) but shifts Wageningen’s Environmental Science lead to the University of Colorado Boulder, confirming that the field-leadership claims should be read as reflecting depth of specialisation rather than unconditional supremacy (Subsection 3.4). OpenAlex’s inconsistent retraction flagging leaves a small, unquantified fraction of h_1 values potentially inflated. The top of the small-institution efficiency table (Table 7) is similarly affected: several entries, including Bellevue University and three other primarily-teaching institutions, owe their apparent efficiency to a handful of author records that are almost certainly disambiguation-merge artifacts rather than genuine researchers (see the discussion following Table 7); this correction was not extended to the rest of the dataset and should be read as a targeted caveat, not a general audit. Saudi Arabia’s efficiency rank (4th, $\varepsilon_3 = 6.69$) is plausibly an artifact of affiliation-listing for funding rather than resident research presence (Bhattacharjee, 2011), a data-quality distortion visible in the output but not removable without ground-truth affiliation data; a multi-institution robustness check (Subsection 3.5) rules out the current-affiliation tiebreak rule specifically as the cause (ε_3 rises rather than falls under that check), but Saudi Arabia is simply the case where this kind of distortion happens to be visible, not necessarily the only country affected by it. Current-affiliation attribution is itself a related, broader limitation: every author’s entire career h_1 is credited to their most recent educational institution, so internationally mobile or multiply affiliated researchers can have career-long impact attributed to wherever they are affiliated now rather than to where the underlying work was produced. A multi-institution robustness check (Subsection 3.5) bounds one specific piece of this problem, the arbitrary tiebreak applied when two or more institutions are equally most-recent, and finds it moves institution-level rankings only modestly ($\rho = 0.966$) and country-level rankings very little ($\rho = 0.997$); it does not establish that current-affiliation attribution more broadly is unproblematic, since it leaves untested every author whose single most-recent affiliation is unambiguous but still not where their cited work was conducted. That broader concern is addressed in this paper by narrowing its claims to what current-affiliation-based h_2 and h_3 can support, not by a further re-attribution that would resolve it. Finally, h_2 is not comparable across fields of different size: Decision Sciences (max $h_2 = 22$) cannot be placed on the same scale as Medicine (max $h_2 = 119$), and the results should be read within fields only.

The most impactful near-term improvement to this framework would require a data source that does not yet exist: departmental affiliation in OpenAlex. If OpenAlex were to record the department within an institution alongside the institution itself, the hierarchy would extend naturally to researcher $\rightarrow$ department $\rightarrow$ institution $\rightarrow$ country, with successive indices h_1 through h_4 . Departments are field-specialized by construction, so the step of inferring each author’s modal field from topic weights may become unnecessary at the departmental level. I flag this as the single most valuable addition OpenAlex could make for scientometric work of this kind.

Statements and Declarations

The author did not receive support from any organization for the submitted work. The author has no relevant financial or non-financial interests to disclose. This study did not involve human participants or animals, and thus no ethical approval or informed consent was required. The data that support the findings of this study are publicly available from the OpenAlex public snapshot. The analysis code is available at the GitHub repository provided in the manuscript.

References

- Alperin, J.P., Portenoy, J., Demes, K., Larivière, V., Haustein, S. (2024). *An analysis of the suitability of OpenAlex for bibliometric analyses*. Retrieved from <https://doi.org/10.48550/arXiv.2404.17663>
- Arencibia-Jorge, R., Barrios-Almaguer, I., Fernández-Hernández, S., Carvajal-Espino, R. (2008). Applying successive H indices in the institutional evaluation: A case study. *Journal of the American Society for Information Science and Technology*, 59(1), 155-157, <https://doi.org/10.1002/asi.20729>
- Arencibia-Jorge, R., & Rousseau, R. (2009). Influence of individual researchers’ visibility on institutional impact: an example of Prathap’s approach to successive h-indices. *Scientometrics*, 79(3), 507–516, <https://doi.org/10.1007/s11192-007-2025-0>
- Bhattacharjee, Y. (2011, December). Saudi universities offer cash in exchange for academic prestige. *Science Magazine*, 334, 1344–1345, <https://doi.org/10.1126/science.334.6061.1344>
- Clauset, A., Shalizi, C.R., Newman, M.E.J. (2009). Power-law distributions in empirical data. *SIAM Review*, 51(4), 661-703, <https://doi.org/10.1137/070710111>

- Culbert, J.H., Hobert, A., Jahn, N., Haupka, N., Schmidt, M., Donner, P., Mayr, P. (2025, April). Reference coverage analysis of OpenAlex compared to Web of Science and Scopus. *Scientometrics*, 130(4), 2475–2492, <https://doi.org/10.1007/s11192-025-05293-3>
- Demes, K. (2026). *A big improvement to our corresponding-author data*. OpenAlex blog. Retrieved from <https://blog.openalex.org/a-big-improvement-to-our-corresponding-author-data/>
- Egghe, L. (2008). Modelling successive h-indices. *Scientometrics*, 77(3), 377–387, <https://doi.org/10.1007/s11192-007-1968-5>
- Hauschke, C., & Nazarovets, S. (2025). (non-) retracted academic papers in OpenAlex. *Journal of Information Science*, 01655515251322478, <https://doi.org/10.1177/01655515251322478>
- Hirsch, J.E. (2005). An index to quantify an individual’s scientific research output. *Proceedings of the National Academy of Sciences*, 102(46), 16569–16572, <https://doi.org/10.1073/pnas.0507655102>
- Mryglod, O., Holovatch, Y., Kenna, R. (2022). Big fish and small ponds: why the departmental h-index should not be used to rank universities. *Scientometrics*, 127(6), 3279–3292, <https://doi.org/10.1007/s11192-022-04373-y>
- Nederhof, A.J. (2006). Bibliometric monitoring of research performance in the Social Sciences and the Humanities: A review. *Scientometrics*, 66(1), 81–100, <https://doi.org/10.1007/s11192-006-0007-2>
- OpenAlex (2024). *OpenAlex: End-to-end process for topic classification*. Google Docs. Retrieved from <https://docs.google.com/document/d/1bDopkhuGieQ4F8gGNj7sEc8WSE8mvLZS/>
- Prathap, G. (2006). Hirsch-type indices for ranking institutions’ scientific research output. *Current Science (00113891)*, 91(11), 1439,
- Priem, J., Piwowar, H., Orr, R. (2022). *OpenAlex: A fully-open index of scholarly works, authors, venues, institutions, and concepts*. Retrieved from <https://doi.org/10.48550/arXiv.2205.01833>

- Ruane, F., & Tol, R. (2008). Rational (successive) h-indices: An application to economics in the Republic of Ireland. *Scientometrics*, 75(2), 395–405, <https://doi.org/10.1007/s11192-007-1869-7>
- Schubert, A. (2007). Successive h-indices. *Scientometrics*, 70(1), 201–205, <https://doi.org/10.1007/s11192-007-0112-x>
- SCImago Research Group (2026). *SCImago institutions rankings 2026*. SCImago Institutions Rankings, Overall Rank export. Retrieved from <https://www.scimagoir.com/getdata.php?ranking=Overall&area=§or=&country=&year=2020&top=08>
- Vaccaro, G., Sánchez-Núñez, P., Witt-Rodríguez, P. (2022). Bibliometrics evaluation of scientific journals and country research output of dental research in Latin America using Scimago Journal and Country Rank. *Publications*, 10(3), <https://doi.org/10.3390/publications10030026>
- van Eck, N.J. (2024, October). *CWTS Leiden ranking open edition 2024 - results*. Zenodo.
- Zhao, R., & Chen, Y. (2025). *Accuracy assessment of OpenAlex and Clarivate Scholar ID with an LLM-assisted benchmark*.
- Zheng, M., Miao, L., Bu, Y., Larivière, V. (2025). Understanding discrepancies in the coverage of OpenAlex: The case of China. *Journal of the Association for Information Science and Technology*, 76(11), 1591-1601, <https://doi.org/https://doi.org/10.1002/asi.70013>